

# COMP2005

---

Derivative Filters

# In 1D

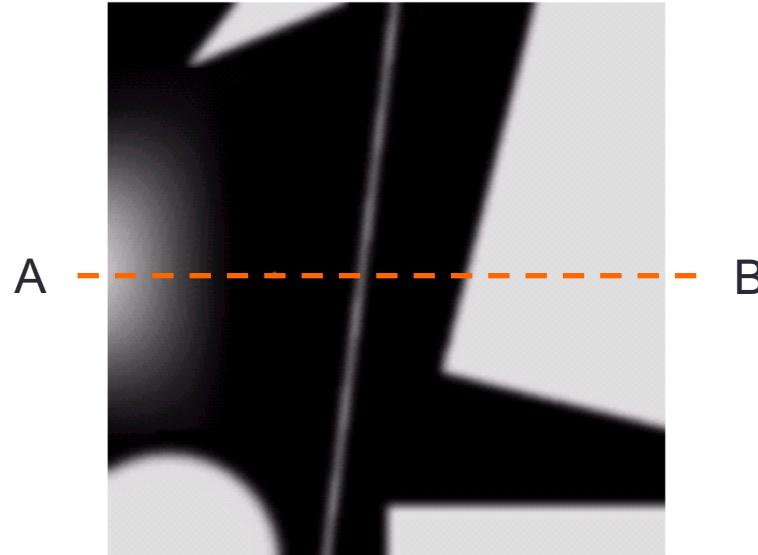
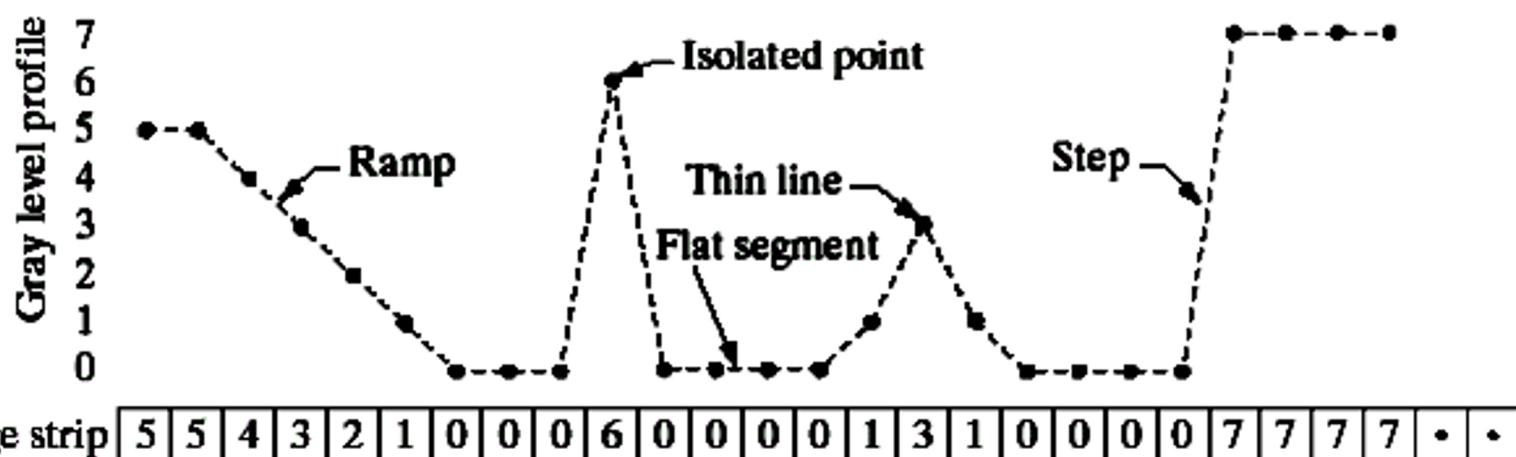


Image features are often characterised by changes in intensity

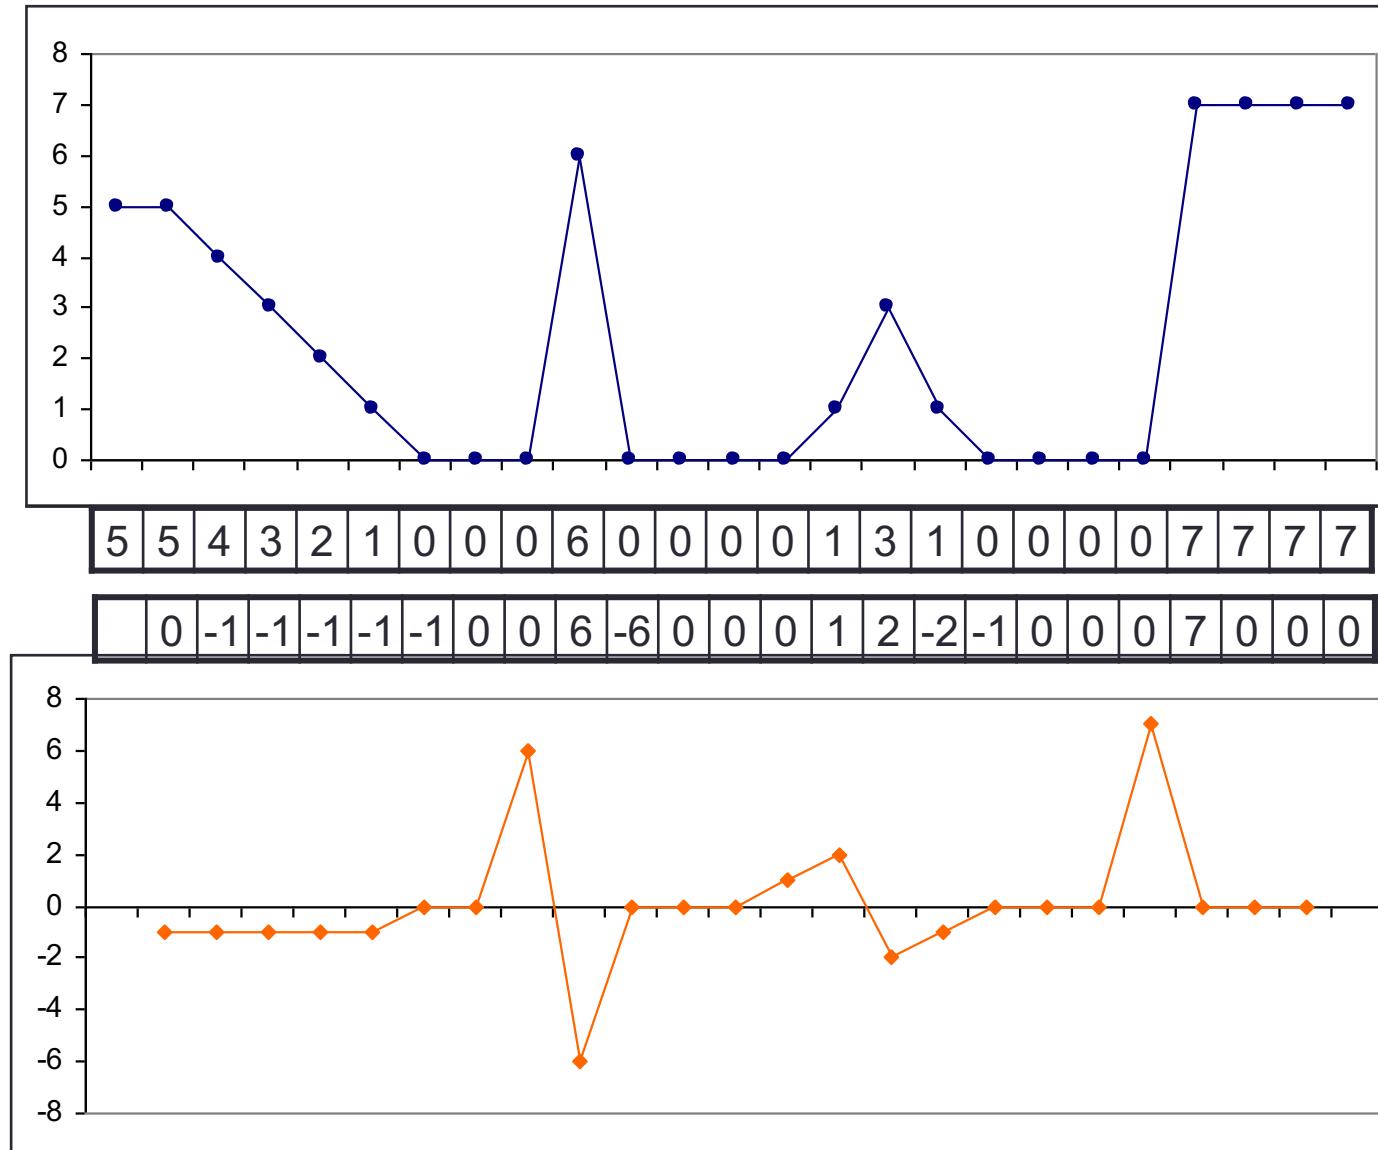


# 1<sup>st</sup> Derivative

- The 1<sup>st</sup> derivative of a function can be approximated by:

$$\frac{\partial f}{\partial x} = f(x+1) - f(x)$$

- The difference between neighbouring values and measures the rate of change of the function
- Can also be approximated by  $(f(x+1) - f(x-1))/2$ , etc.



## Raw data

# 1st derivative

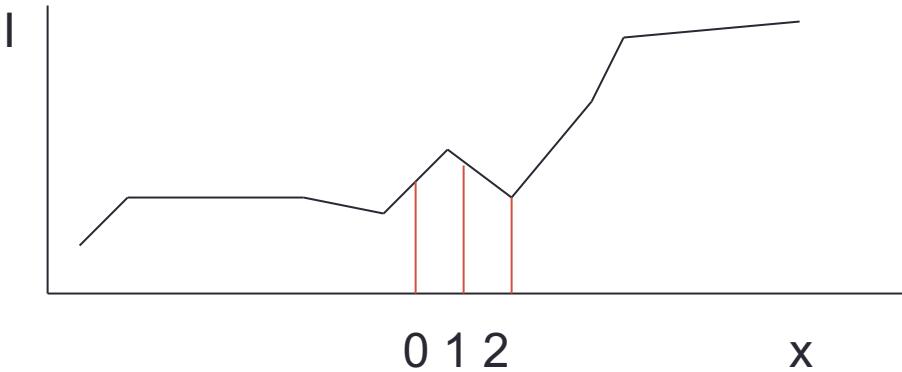
## 2<sup>nd</sup> Derivative

- The formula for the 2<sup>nd</sup> derivative of a function is:

$$\frac{\partial^2 f}{\partial^2 x} = f(x+1) + f(x-1) - 2f(x)$$

- Simply takes into account the values both before and after the current value
- Derived by estimating the 1<sup>st</sup> derivative at  $x + 0.5$  and  $x - 0.5$  and computing the derivative of the resulting data

## 2<sup>nd</sup> Derivative

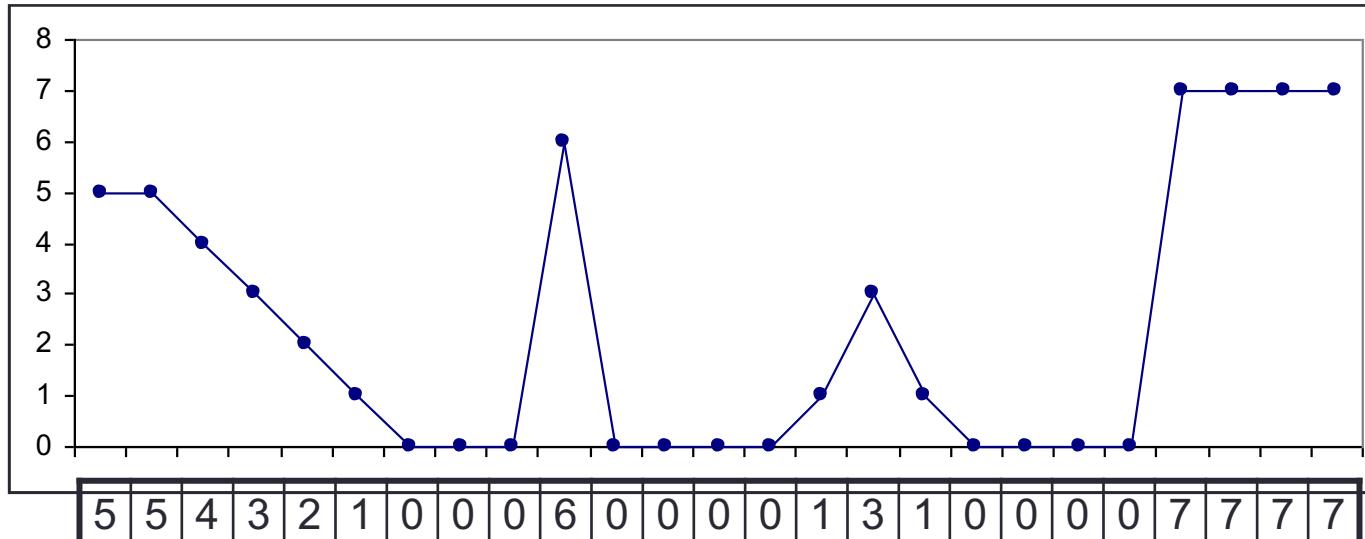


$$I''(1) = (I'(1.5) - I'(0.5))/1$$

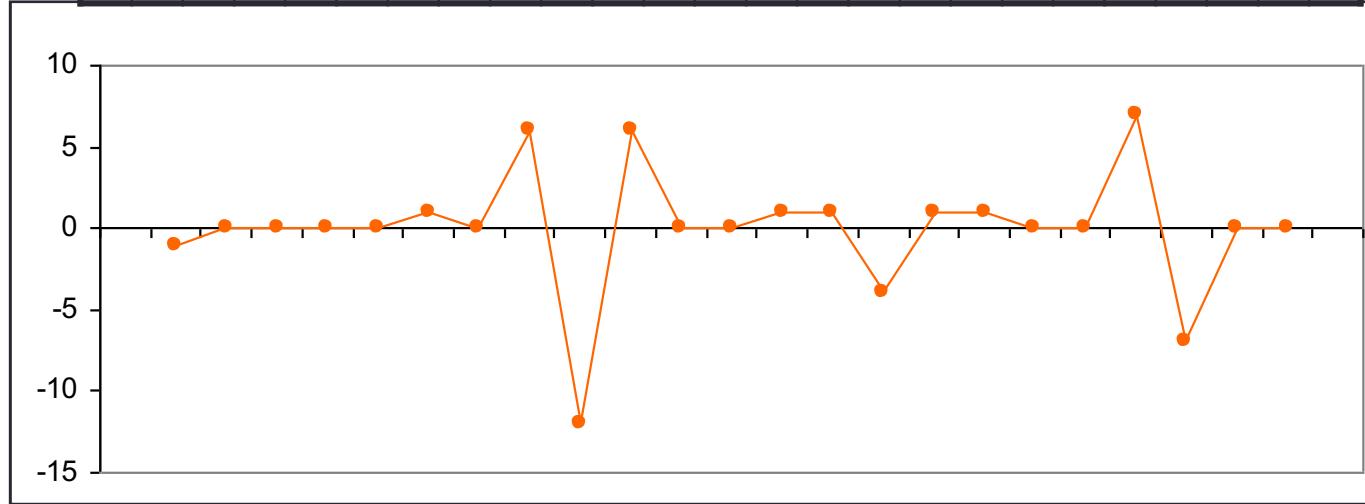
$$I'(0.5) = (I(1) - I(0))/1 \quad \text{and} \quad I'(1.5) = (I(2) - I(1))/1$$

$$\therefore I''(1) = 1 \cdot I(0) - 2 \cdot I(1) + 1 \cdot I(2)$$

1	-2	1
---	----	---



Raw data



2<sup>nd</sup>  
derivative

# Derivatives in 2D

- 2<sup>nd</sup> derivatives generalize to 2D quite easily, implementing a 1<sup>st</sup> derivative in 2D is a little more complex
- For a function  $f(x, y)$  the gradient of  $f$  at coordinates  $(x, y)$  is given as the column vector:

$$\nabla f = \begin{bmatrix} G_x \\ G_y \end{bmatrix} = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix}$$

- Computation of the 1<sup>st</sup> derivative can't be done by convolution alone

# 1<sup>st</sup> Derivative Filtering

- The magnitude of the 1<sup>st</sup> derivative vector is

$$\begin{aligned}\nabla f &= \text{mag}(\nabla f) \\ &= [G_x^2 + G_y^2]^{1/2} \\ &= \left[ \left( \frac{\partial f}{\partial x} \right)^2 + \left( \frac{\partial f}{\partial y} \right)^2 \right]^{1/2}\end{aligned}$$

which can be simplified to

$$\nabla f \approx |G_x| + |G_y|$$

# 1<sup>st</sup> Derivative Filters

- Many 1<sup>st</sup> derivative filters have been proposed

- Roberts' Cross Operators

1	0
0	-1

0	1
-1	0

- Sobel Operators

-1	0	1
-2	0	2
-1	0	1

$G_x$

-1	-2	-1
0	0	0
1	2	1

$G_y$

These operators are most commonly associated with edge detection

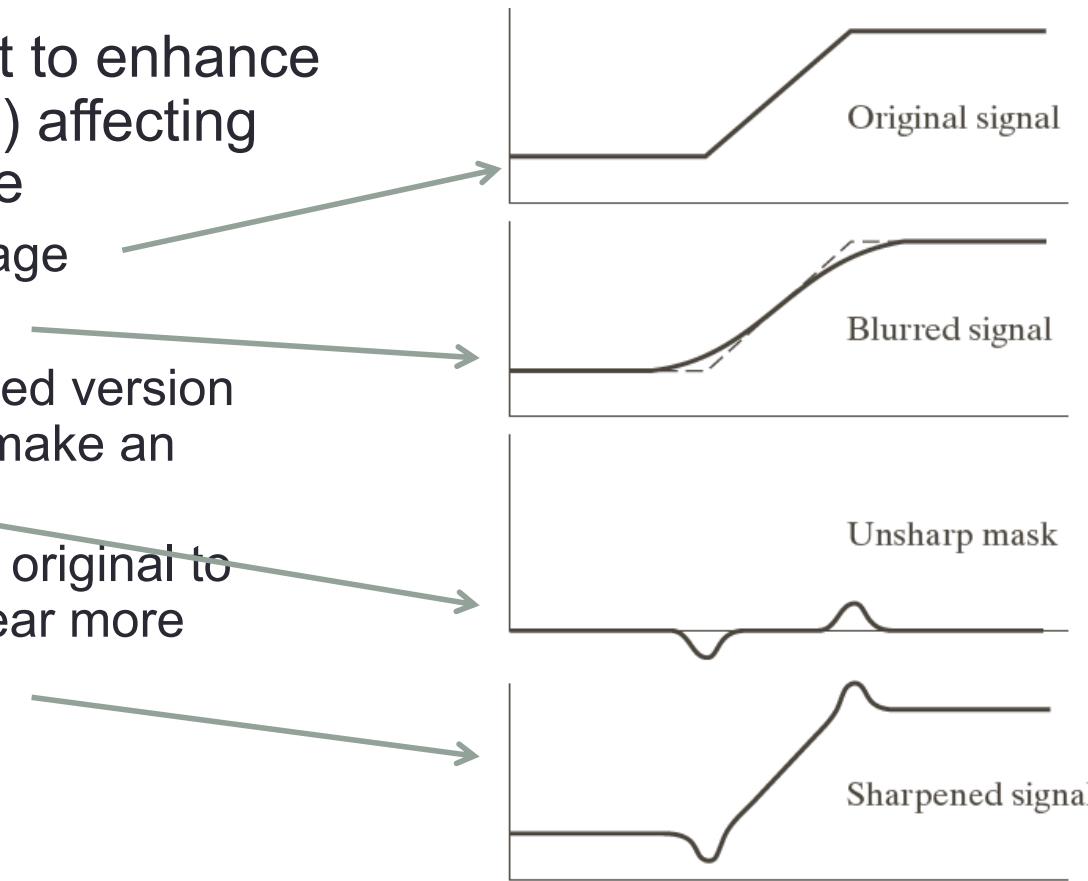
# COMP2005

---

Image Sharpening

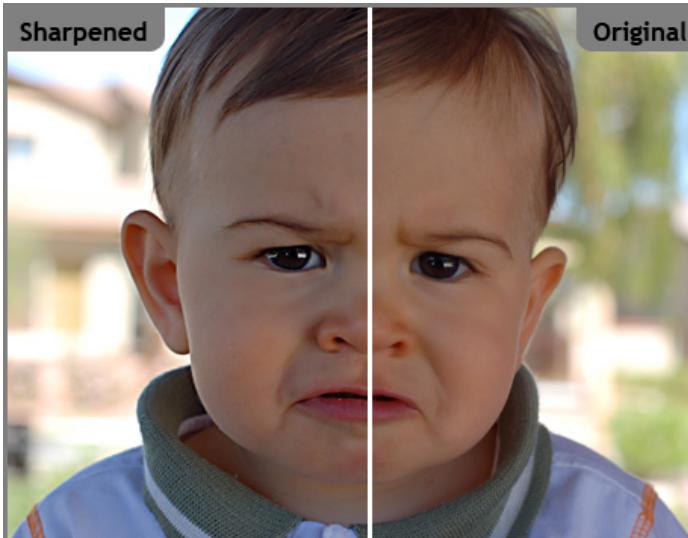
# Edge Enhancement: Unsharp Masking

- Edges are important
- Sometimes we want to enhance them without (much) affecting the rest of the image
  - Take the original image
  - Gaussian smooth it
  - Subtract the smoothed version from the original to make an *unsharp mask*
  - Add the mask to the original to make the edge appear more obvious



# Unsharp Masking

- Makes edges noticeably sharper
  - Even if they are noise
  - Sometimes too much



# Derivative Filters

- Unsharp filtering enhances edges by comparing the original with a smoothed image
  - relies on the smoothing effect of a Gaussian function introducing a difference between original and processed images
  - parameterised by  $\sigma$
  - simple, but effect is hard to predict, so hard to parameterise
- A more direct way to highlight edges and other features associated with high image gradients is to estimate derivatives.....

# Image Sharpening with Derivatives

- The 2<sup>nd</sup> derivative is more useful for image enhancement than the 1<sup>st</sup> derivative
  - Stronger response to fine detail
  - Simpler implementation
- The most common sharpening filter is the *Laplacian*
  - Isotropic
  - One of the simplest sharpening filters
  - Straightforward digital implementation via convolution

# The Laplacian

$$\nabla^2 f = \frac{\partial^2 f}{\partial^2 x} + \frac{\partial^2 f}{\partial^2 y}$$

$$\frac{\partial^2 f}{\partial^2 x} = f(x+1, y) + f(x-1, y) - 2f(x, y)$$

$$\frac{\partial^2 f}{\partial^2 y} = f(x, y+1) + f(x, y-1) - 2f(x, y)$$

# The Laplacian

$$\begin{aligned}\nabla^2 f = & [f(x+1, y) + f(x-1, y) \\& + f(x, y+1) + f(x, y-1)] \\& - 4f(x, y)\end{aligned}$$

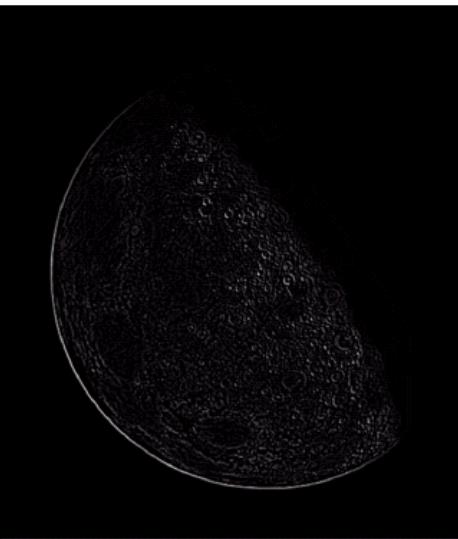
0	1	0
1	-4	1
0	1	0

# The Laplacian

- Highlights edges and other discontinuities



Original  
Image



Laplacian  
Filtered Image



Laplacian  
Filtered Image  
Scaled for Display

# Could do better?

- The result of Laplacian filtering is not an enhanced image
  - subtract the Laplacian result from the original image to generate the final sharpened enhanced image,  
c.f. unsharp masking

$$g(x, y) = f(x, y) - \nabla^2 f$$



Original



Laplacian Filtered

=



Sharpened

# Laplacian Enhancement

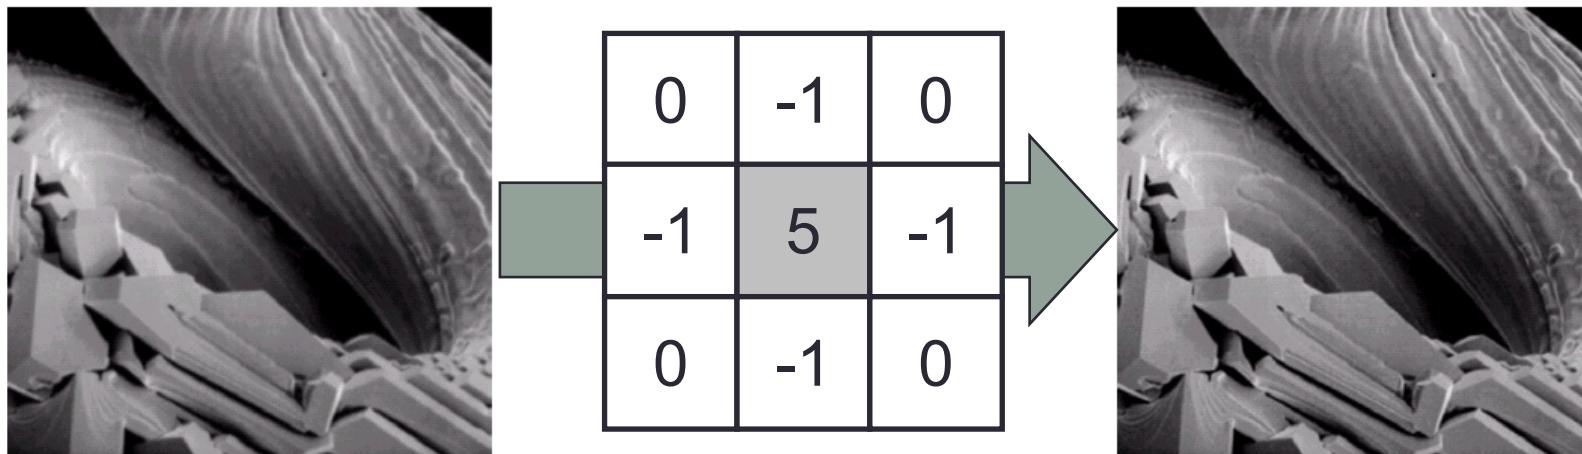


# A Single Enhancement Operator

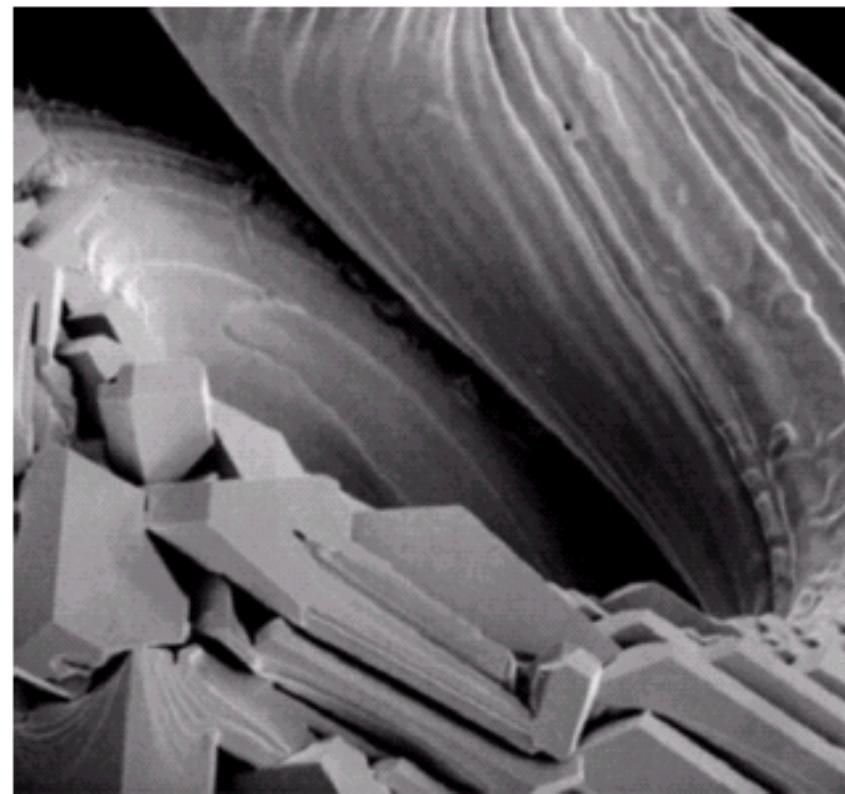
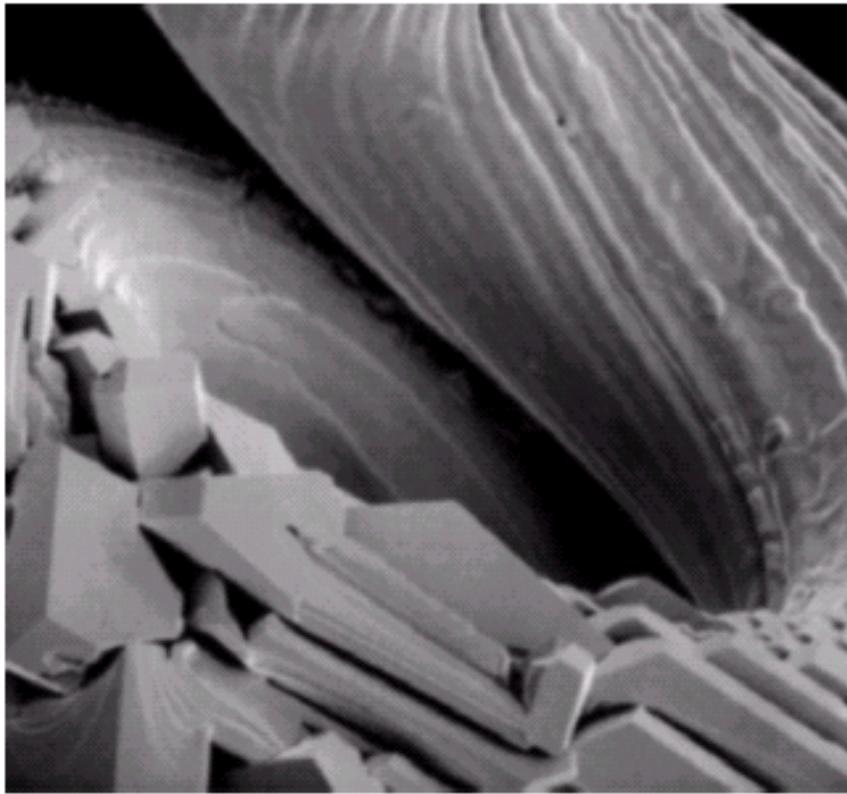
$$\begin{aligned}g(x, y) &= f(x, y) - \nabla^2 f \\&= f(x, y) - [f(x+1, y) + f(x-1, y) \\&\quad + f(x, y+1) + f(x, y-1) \\&\quad - 4f(x, y)] \\&= 5f(x, y) - f(x+1, y) - f(x-1, y) \\&\quad - f(x, y+1) - f(x, y-1)\end{aligned}$$

# A Single Operator

- Convolution with this operator performs image sharpening in a single step



# A Single Operator



# Variations on the Theme

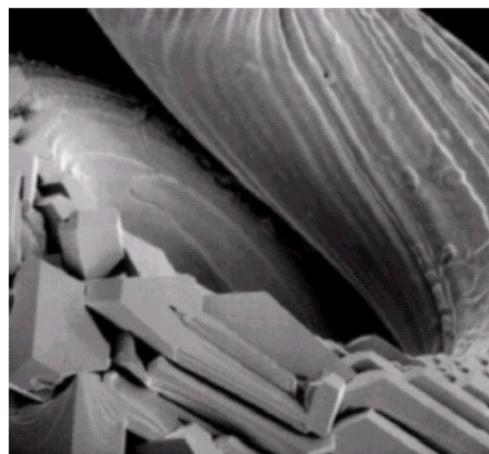
- Some formulations take 2<sup>nd</sup> derivatives measured across the diagonals into account

0	1	0
1	-4	1
0	1	0

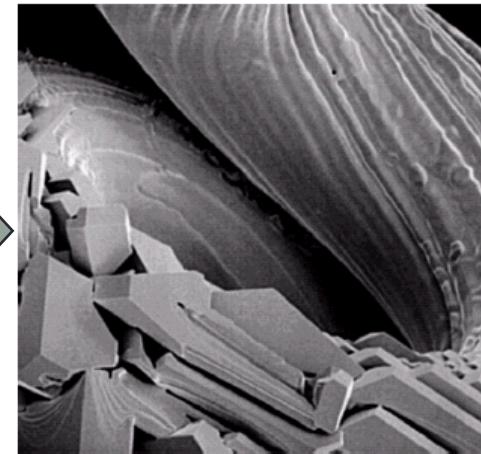
Basic

1	1	1
1	-8	1
1	1	1

Extended



-1	-1	-1
-1	9	-1
-1	-1	-1



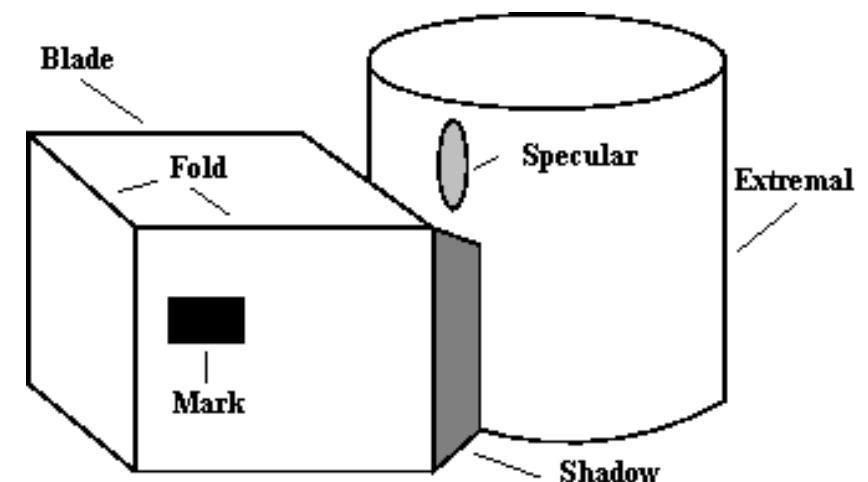
# COMP2005

---

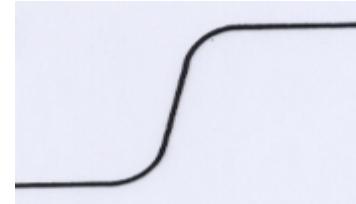
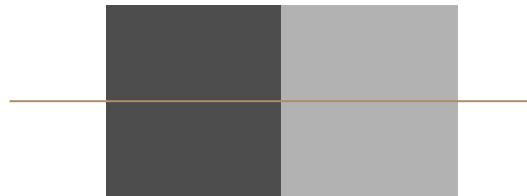
What is Edge Detection?

# Edge Detection

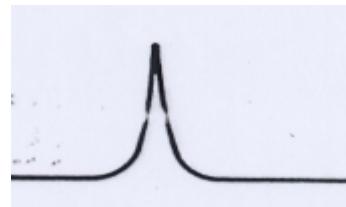
- First step in many image analysis and computer vision processes and applications
- The goal is to mark points at which image intensity changes sharply – edges
- Sharp changes in **image** properties reflect events/changes in the **world**
- This is only an assumption, but it is usually true



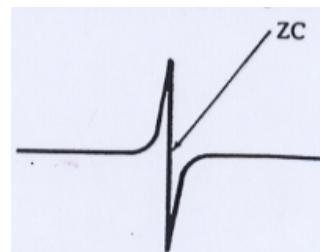
# The Theory



- An idealised image of an edge



- 1<sup>st</sup> derivative: a peak



- 2<sup>nd</sup> derivative: a zero-crossing

To detect edges find peaks in the 1<sup>st</sup> derivative of intensity or zero-crossings in the 2<sup>nd</sup> derivative

# The Result

```
>> im = imread('cameraman.tif');
>> edges = edge(im, 'Canny');
>> imshowpair(im, edges, 'montage');
```



# COMP2005

---

Edge Detection using 1<sup>st</sup> Derivative Filters

# 1<sup>st</sup> Derivative Filters

- Roberts' Cross Operators

1	0
0	-1

0	1
-1	0

- Sobel Operators

-1	0	1
-2	0	2
-1	0	1

-1	-2	-1
0	0	0
1	2	1

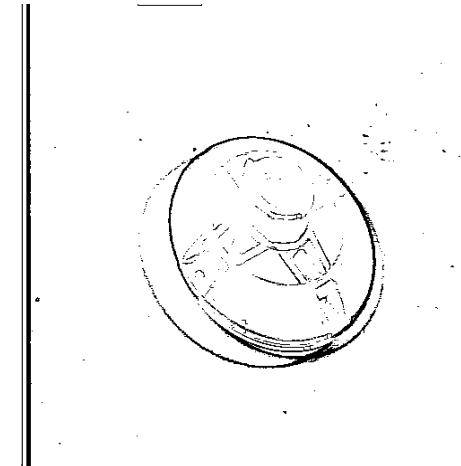
- Applied separately and results combined to estimate magnitude

# Detection & Thresholding

- Significant peaks in magnitude of 1<sup>st</sup> derivative are high
- Apply a threshold, all peaks higher than the threshold value are significant, all others are ignored



■ Too low



■ Too high

# Edge Magnitude & Direction

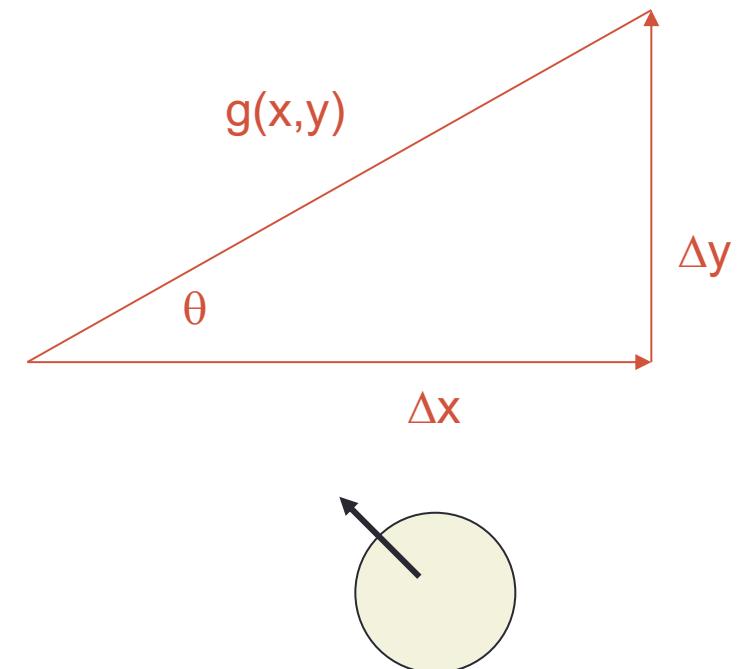
- For an image function,  $I(x,y)$ , the gradient magnitude,  $g(x,y)$  is given by

$$g(x,y) \cong (\Delta x^2 + \Delta y^2)^{1/2}$$

- The gradient direction,  $\theta(x,y)$ , gives the direction of steepest image gradient

$$\theta(x,y) \cong \text{atan}(\Delta y / \Delta x)$$

- This gives the direction of a line perpendicular to the edge

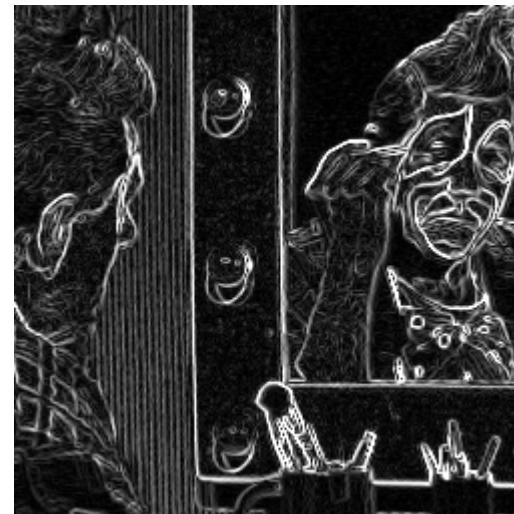


# Roberts' Cross Operator

- Very quick to compute - 4 pixels, only subtractions and additions, but is very sensitive to noise and only gives a strong response to very sharp edges



Original



Cross Operator



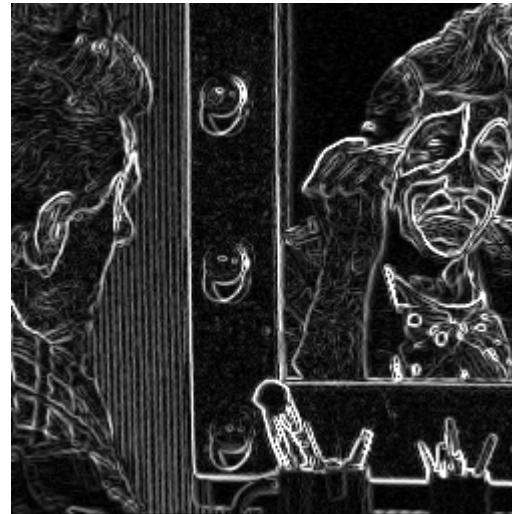
Thresholded

# Sobel vs Roberts

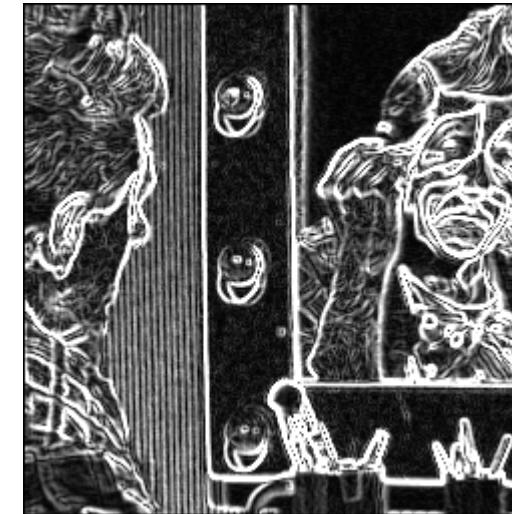
- Both use a user-supplied threshold. Sobel is still in use, Roberts is less common.
- Larger Sobel operators are more stable in noise



Original



Roberts



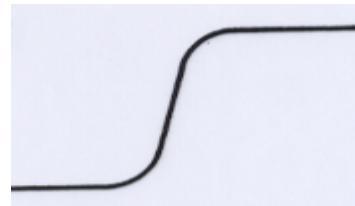
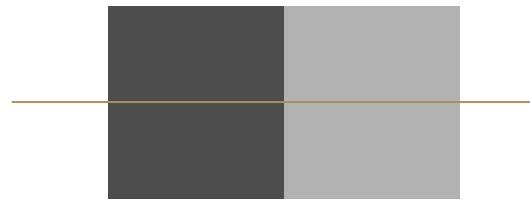
Sobel

# COMP2005

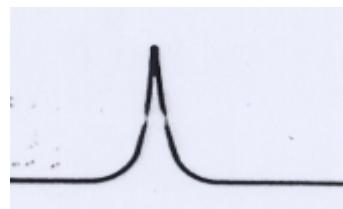
---

Edge Detection using 2<sup>nd</sup> Derivatives

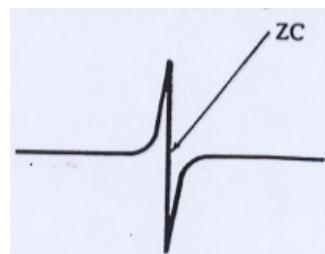
# The Theory



- An idealised image of an edge



- 1<sup>st</sup> derivative: a peak



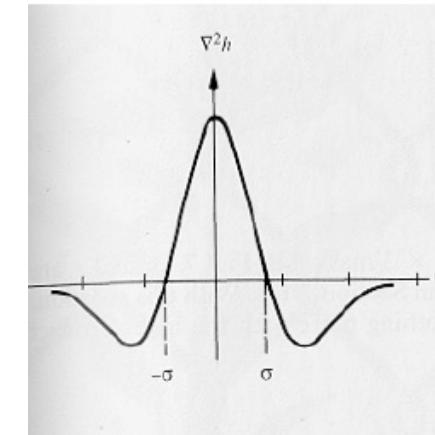
- 2<sup>nd</sup> derivative: a zero-crossing

To detect edges find peaks in the 1<sup>st</sup> derivative of intensity or zero-crossings in the 2<sup>nd</sup> derivative

# 2<sup>nd</sup> Derivatives: Marr-Hildreth

- Biologically inspired
  - Gaussian smooth, compute Laplacian
- OR
- Convolve with the **Laplacian of a Gaussian**

$$\nabla^2[f(x, y) * G(x, y)] = \nabla^2G(x, y) * f(x, y)$$

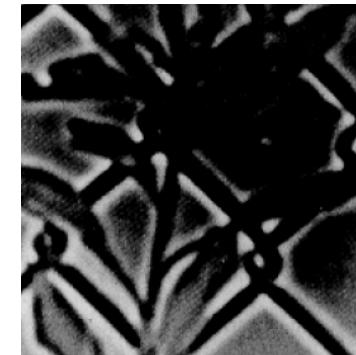


# Laplacian of Gaussian (LoG)

5 × 5 Laplacian of Gaussian mask

$$\begin{bmatrix} 0 & 0 & -1 & 0 & 0 \\ 0 & -1 & -2 & -1 & 0 \\ -1 & -2 & 16 & -2 & -1 \\ 0 & -1 & -2 & -1 & 0 \\ 0 & 0 & -1 & 0 & 0 \end{bmatrix}$$

17 × 17 Laplacian of Gaussian mask																
0	0	0	0	0	0	-1	-1	-1	-1	0	0	0	0	0	0	0
0	0	0	0	-1	-1	-1	-1	-1	-1	0	0	0	0	0	0	0
0	0	-1	-1	-1	-2	-3	-3	-3	-3	-2	-1	-1	0	0	0	0
0	0	-1	-1	-2	-3	-3	-3	-3	-3	-3	-2	-1	0	0	0	0
0	-1	-1	-2	-3	-3	-3	-2	-3	-2	-3	-3	-2	-1	-1	0	0
0	-1	-2	-3	-3	-3	0	2	4	2	0	-3	-3	-3	-2	-1	0
-1	-1	-3	-3	-3	0	4	10	12	10	4	0	-3	-3	-3	-1	-1
-1	-1	-3	-3	-2	2	10	18	21	18	10	2	-2	-3	-3	-1	-1
-1	-1	-3	-3	-3	4	12	21	24	21	12	4	-3	-3	-3	-1	-1
-1	-1	-3	-3	-2	2	10	18	21	18	10	2	-2	-3	-3	-1	-1
-1	-1	-3	-3	-3	0	4	10	12	10	4	0	-3	-3	-3	-1	-1
0	-1	-2	-3	-3	0	2	4	2	0	-3	-3	-3	-2	-1	0	0
0	-1	-1	-2	-3	-3	-2	-3	-2	-3	-3	-2	-1	-1	0	0	0
0	0	-1	-1	-2	-3	-3	-3	-3	-3	-3	-2	-1	-1	0	0	0
0	0	-1	-1	-1	-2	-3	-3	-3	-3	-2	-1	-1	-1	0	0	0
0	0	0	0	-1	-1	-1	-1	-1	-1	0	0	0	0	0	0	0



filtering



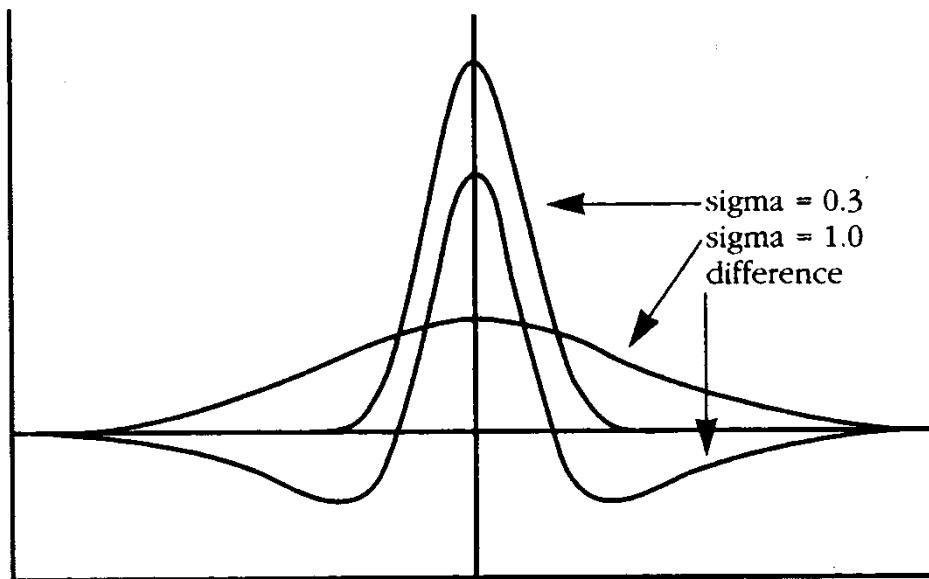
zero-crossings

# Difference of Gaussians

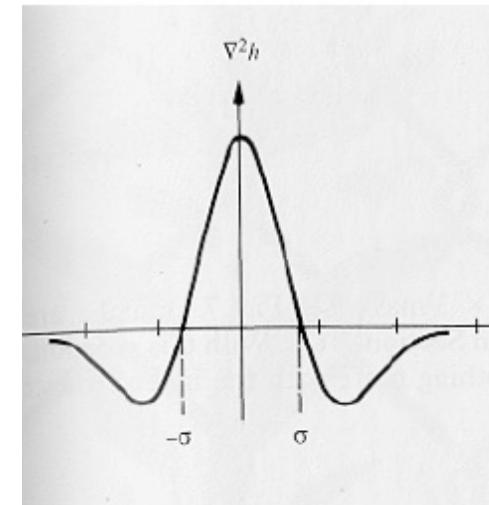
- The Laplacian of Gaussian can be approximated by the difference between two Gaussian functions:

$$\nabla^2 G \approx G(x, y; \sigma_1) - G(x, y; \sigma_2)$$

approximation



actual LoG



# DoG Filtering

$$\nabla^2 G \approx G(x, y; \sigma_1) - G(x, y; \sigma_2)$$



$\sigma = 1$



$\sigma = 2$

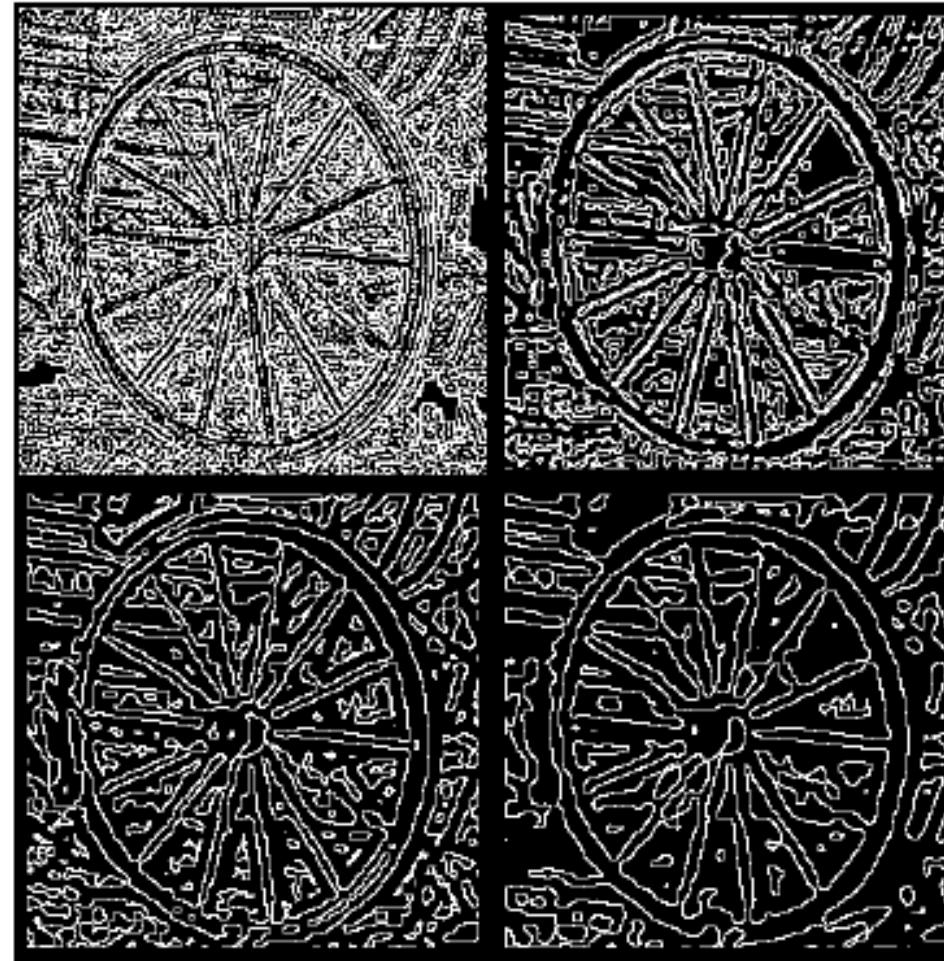


(b)-(a)

Ratio ( $\sigma_1/\sigma_2$ ) for best approximation is about 1.6.  
(Some people like  $\sqrt{2}$ .)

# Marr-Hildreth

- Choice of  $\sigma$  gives flexibility



$\sigma = 1$	$\sigma = 2$
$\sigma = 3$	$\sigma = 4$

# 1<sup>st</sup> vs 2<sup>nd</sup> Derivative Methods

- Peaks in 1<sup>st</sup> derivative
  - strong response at edges, but also respond to noise
  - Peak detection and threshold selection need care
- Zero crossings in 2<sup>nd</sup> derivative
  - well-defined, easy to detect
  - must form smooth, connected contours
  - tend to round off corners
- 1<sup>st</sup> derivative methods are much more common in practical applications,
  - in part because of John Canny

# COMP2005

---

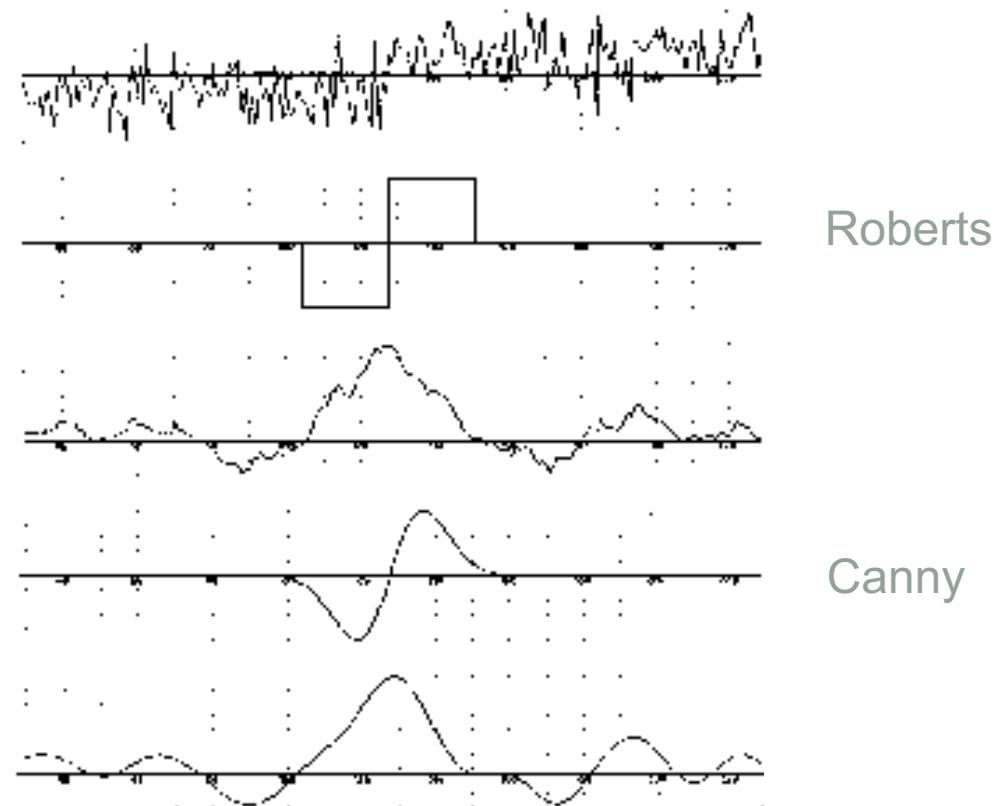
The Canny Operator

# What Canny Did

- John Canny tried to find the optimal edge detector, assuming a perfect step edge in Gaussian noise
- Optimal means
  - **Good Detection** - it should mark all the edges and only all the edges
  - **Good Localisation** - the points marked should be as close to the real edge as possible
  - **Minimal Response** - each edge should be reported only once
- Canny used the *Calculus of Variations*: finds the function which best satisfies some functional

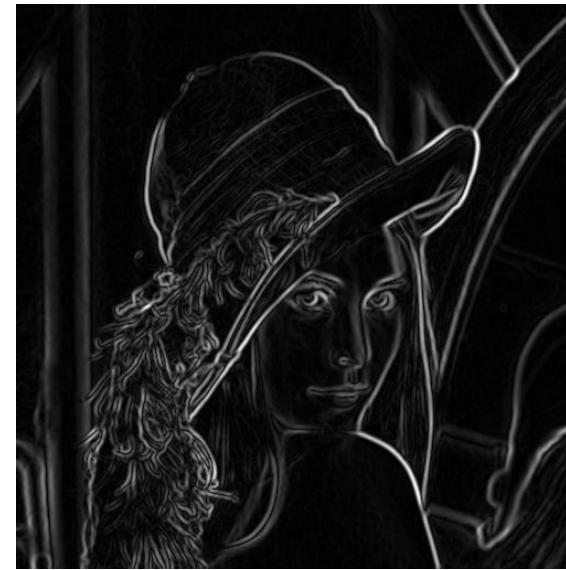
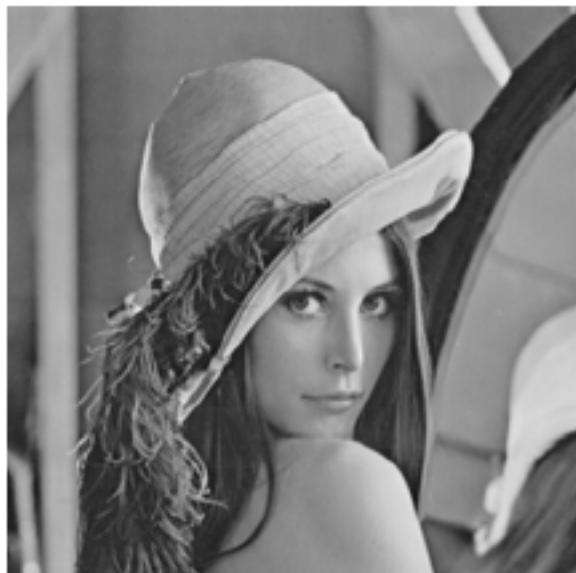
# The Canny Operator

- The optimal detector was a sum of 4 exponential terms, but is very closely approximated by the 1st derivative of a Gaussian
  - i.e. 1<sup>st</sup> derivative of a Gaussian smoothed image
- Gives a cleaner response to a noisy edge than square operators
- Most implementations are 2D Gaussian smoothing + Roberts style derivative



# Non-Maximal Suppression

- The Canny operator's response is cleaner than Sobel or Roberts, but it needs an explicit step to enforce Minimal Response



Canny Operator

# Non-Maximal Suppression

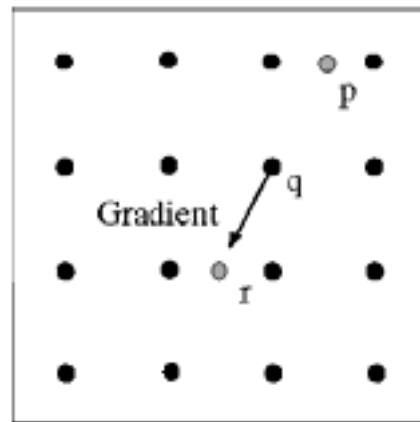
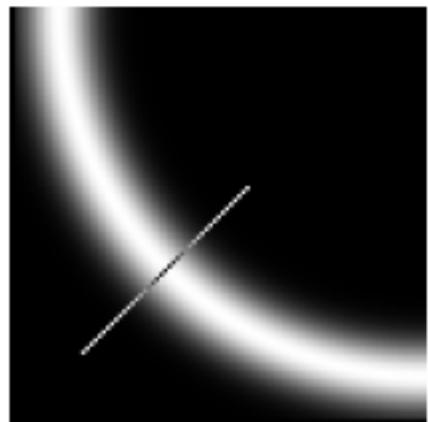
- Thresholding raw operator response would leave thick lines



How to turn  
these thick  
regions of the  
gradient into  
curves?

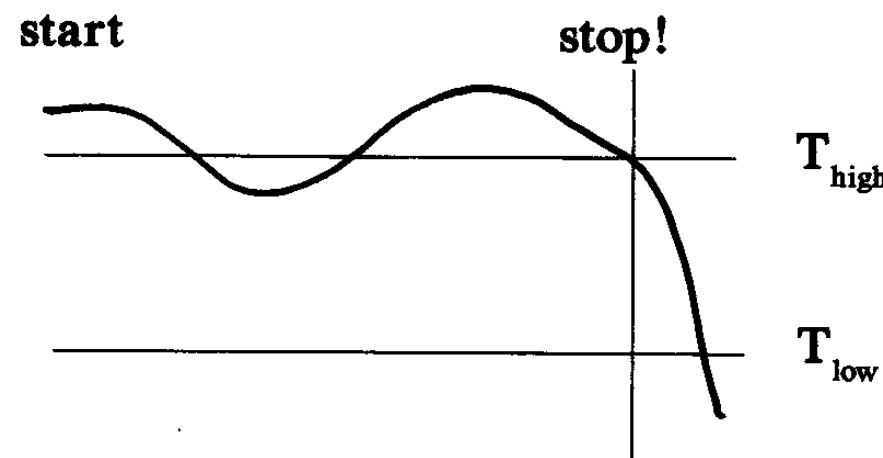
# Non-Maximal Suppression

- Check if pixel is a local maximum along the gradient direction
- Select a single maximum across the width of the edge



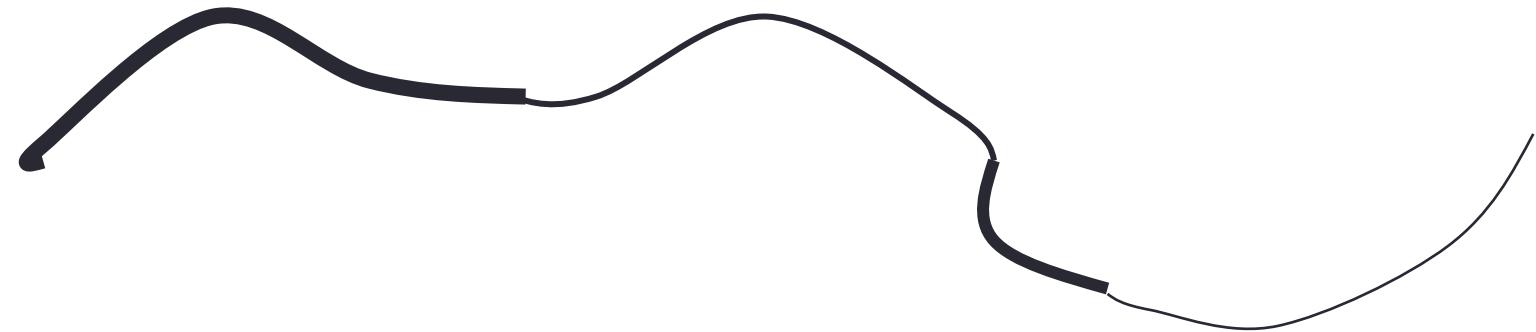
# Thresholding with Hysteresis

- Simple thresholding tests each pixel independently: edges aren't really independent, they make up lines
- The industry standard edge thresholding method
  - Allows a band of variation, but assumes continuous edges
  - User still selects parameters, but it's easier, less precise



# Thresholding with Hysteresis

- The effect is to keep weak edges if they connect strong edges, as long as
  - the strong edges are really strong
  - the weak edges aren't really weak



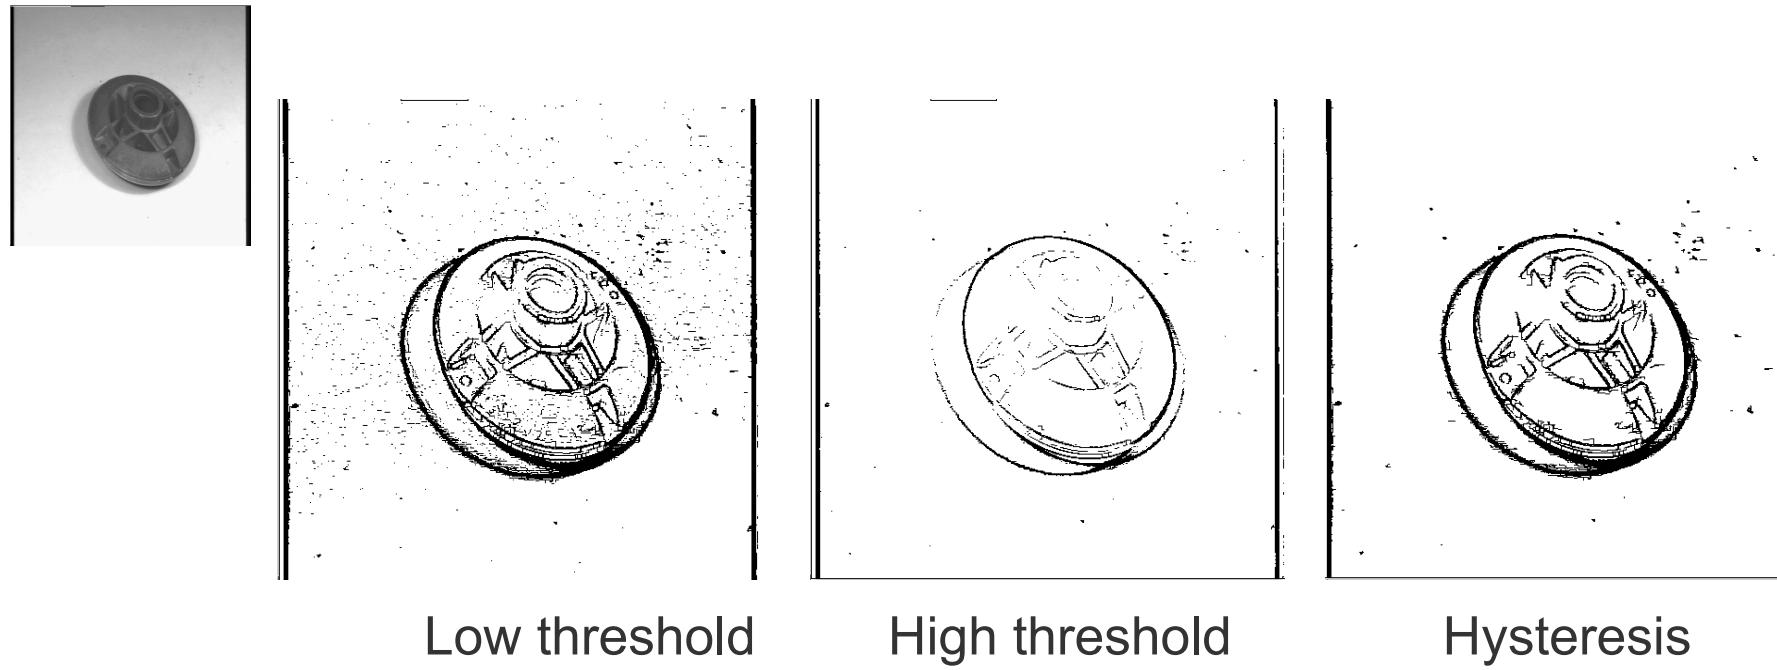
# Thresholding with Hysteresis

Hysteresis fills in most of most of the gaps



Problem:  
pixels along  
this edge  
didn't  
survive the  
thresholding

# Thresholding with Hysteresis



# What Canny Did

- Showed that 1<sup>st</sup> derivative of a Gaussian smoothed image is the optimal way to detect step edges in noise
  - Explained why 1<sup>st</sup> derivatives are a good idea
- Designed the industry standard thresholding method
  - Non-maximal suppression
  - Thresholding with hysteresis
- Effectively solved the edge detection problem

# COMP2005

---

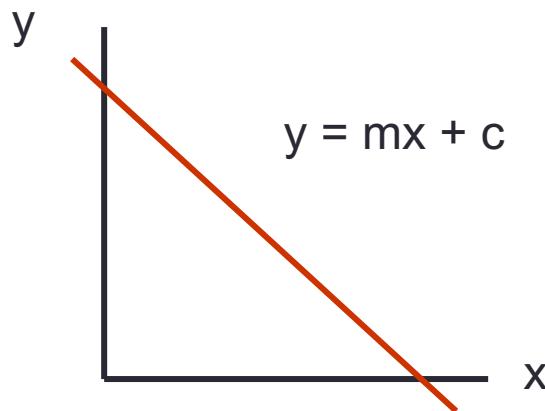
The Hough Transform

# Line Parameters

- The standard equation of a line is

$$y = mx + c$$

- If we know the line parameters  $m, c$  we can vary  $x$ , compute  $y$  and draw the line
- This line represents the set of  $(x, y)$  pairs that satisfy the above equation



But we don't know the line **parameters** - we just have some data points  $(x_i, y_i)$

# Parameter Space

- But if a data point  $(x_i, y_i)$  lies on a line then that line's parameters  $m, c$  must satisfy

$$y_i = mx_i + c$$

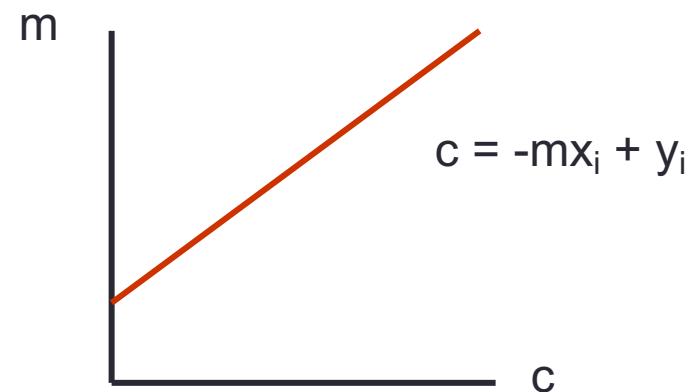
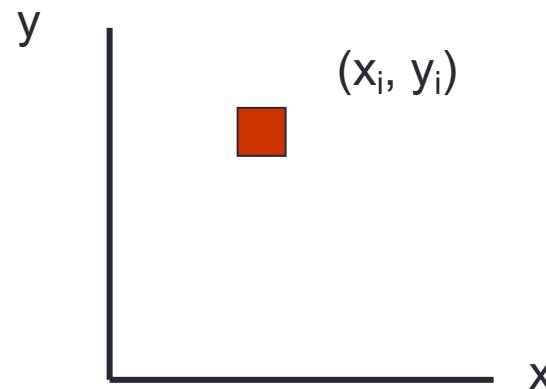
- So we can represent all possible lines through  $(x_i, y_i)$  by the set of  $(m, c)$  pairs that satisfy this equation

- Rearranging the equation gives

$$c = -mx_i + y_i$$

- This also describes a straight line, but as  $x_i$  and  $y_i$  are known and  $m$  and  $c$  are unknown that line is in  $m, c$  space – a parameter space

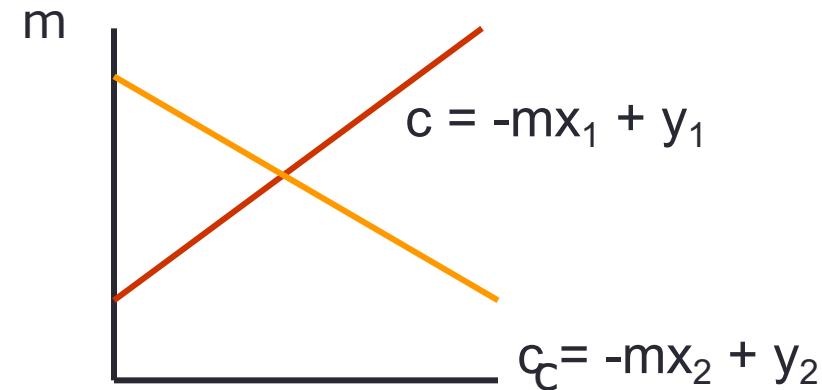
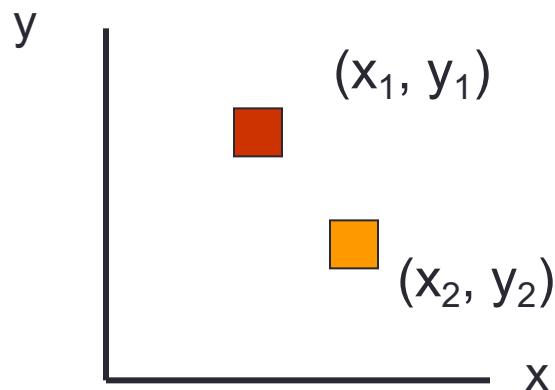
# Parameter Space



- This straight line in  $m,c$  space represents all the values of  $(m,c)$  that satisfy  $y_i = mx_i + c$ , and so **the parameters of all the lines that could pass through  $(x_i, y_i)$**

# Parameter Space

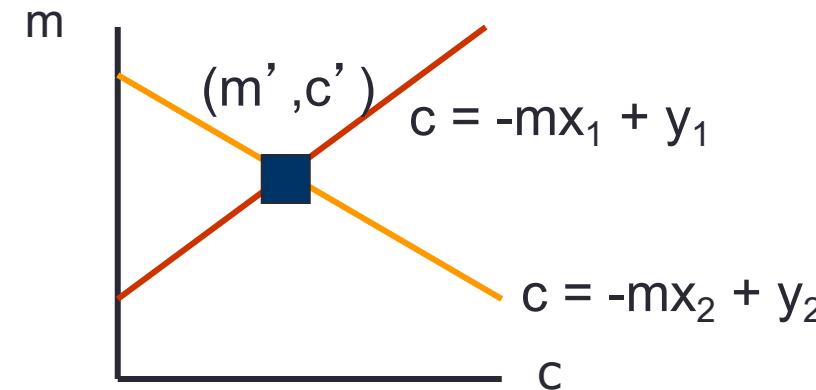
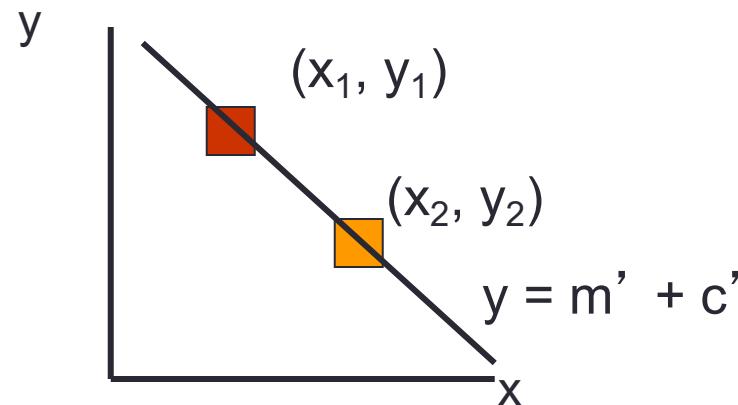
- Suppose there are 2 edges in our image  $(x_1, y_1)$  and  $(x_2, y_2)$



- They will each generate a line in  $(m, c)$  space representing the set of lines that could pass through them

# Parameter Space

- The intersection of our 2 lines is the intersection of 2 sets of line parameters



- The point of intersection therefore gives the parameters of a line through both  $(x_1, y_1)$  and  $(x_2, y_2)$

# Parameter Space

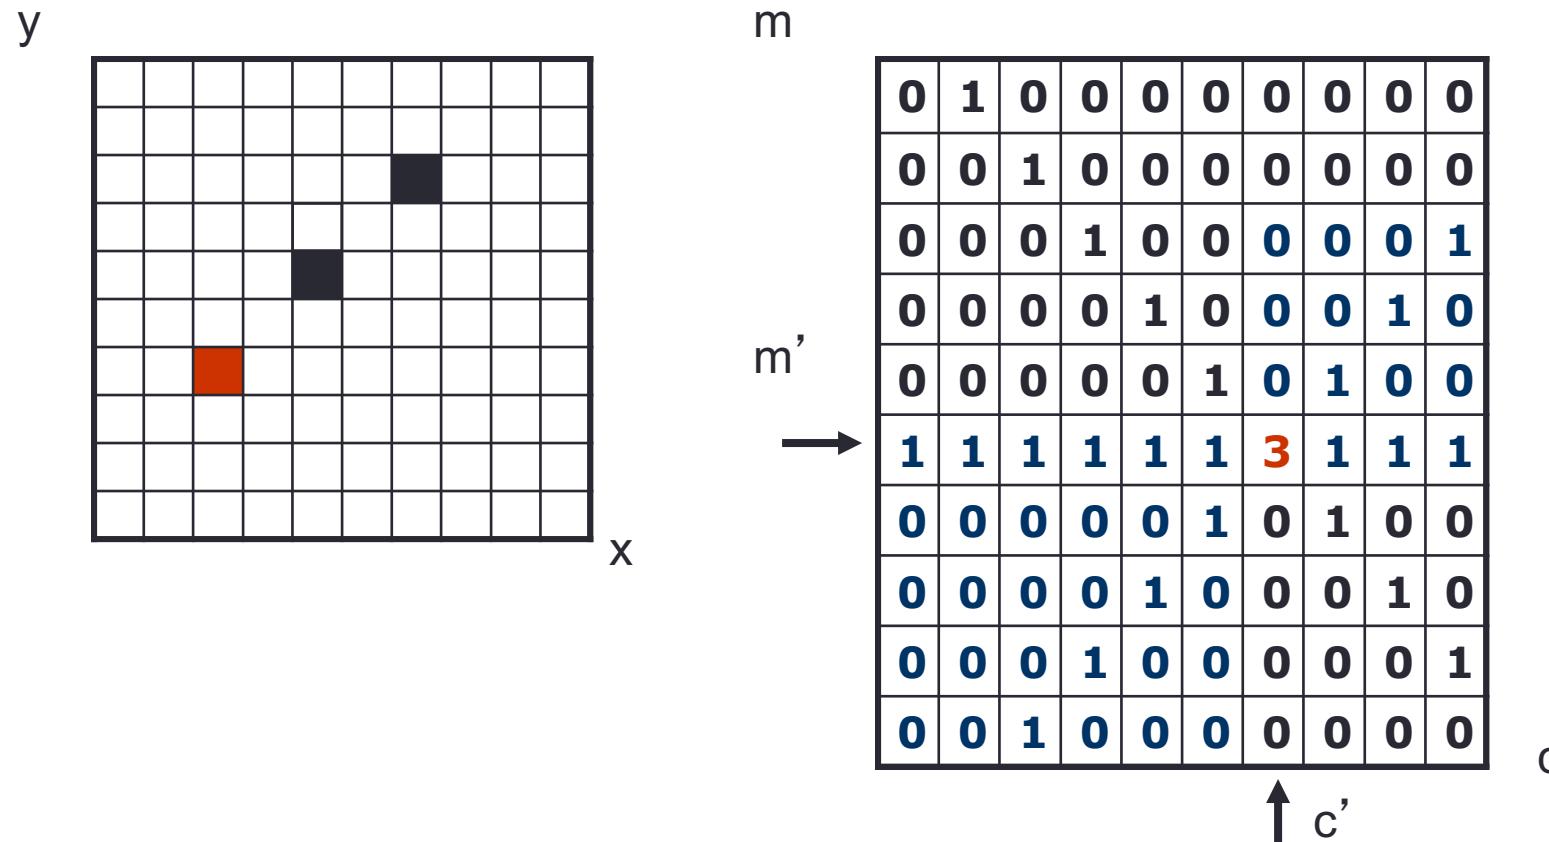
- Taking this one step further
  - All pixels which lie on a line in  $(x,y)$  space are represented by lines which all pass through a single point in  $(m,c)$  space
  - The point through which they all pass gives the values of  $m'$  and  $c'$  in the equation of the line:  $y=m'x+c'$ .
- To detect lines all we need to do is transform each edge point into  $m,c$  space and look for places where lots of lines intersect

*This is what Hough Transforms do*

# The Hough Transform

1. Quantise  $(m,c)$  space into a two-dimensional array  $A$  for appropriate steps of  $m$  and  $c$
2. Initialise all elements of  $A(m,c)$  to zero
3. For each pixel  $(x_i, y_i)$  which lies on some edge in the image, we add 1 to all elements of  $A(m,c)$  whose indices  $m$  and  $c$  satisfy  $y_i = mx_i + c$
4. Search for elements of  $A(m,c)$  which have large values - each one found corresponds to a line in the original image

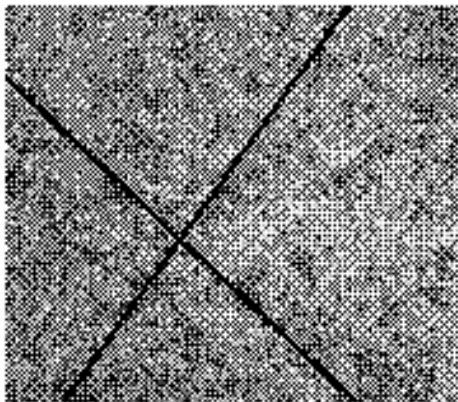
# The Hough Transform



There is one line:  $y = m'x + c'$

# A (More) Real Example

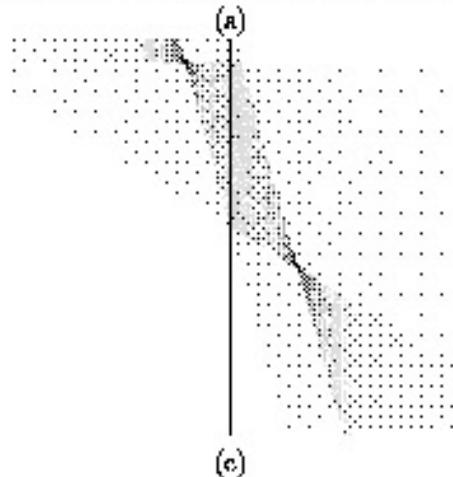
Image



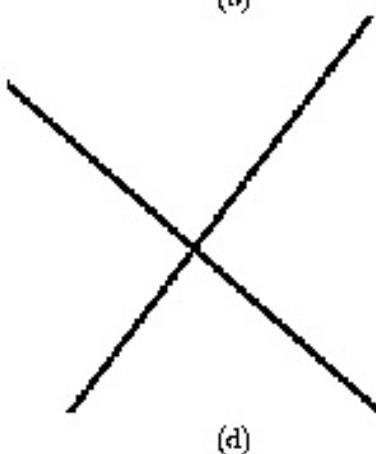
Edges



Accumulators

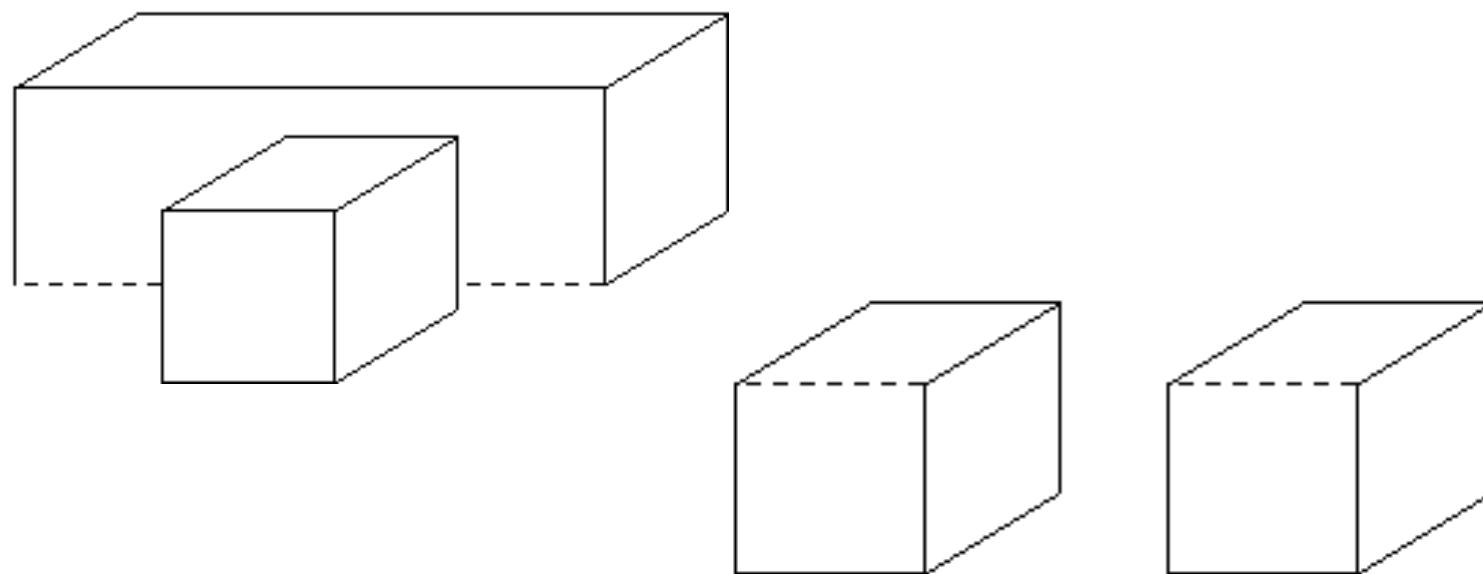


Result



# Line Parameters

- The Hough Transform only supplies the parameters of the lines it detects; this can be an advantage or a disadvantage



# COMP2005

---

Other Parameter Spaces

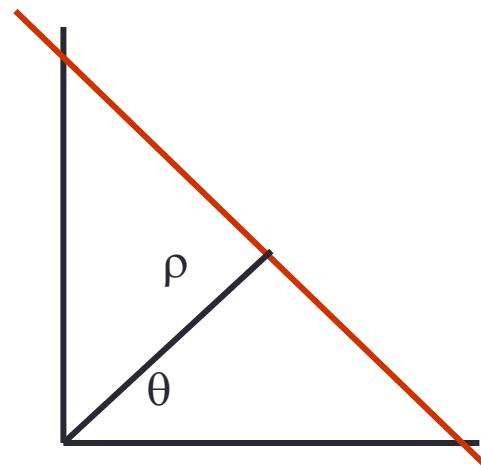
# Straight Lines Again

- Another equation for a straight line

$$\rho = x \cdot \cos(\theta) + y \cdot \sin(\theta)$$

- For an  $n \times n$  image

- $\rho$  is in range  $[0, 2n]$
  - $\theta$  is in range  $[0, 2\pi]$

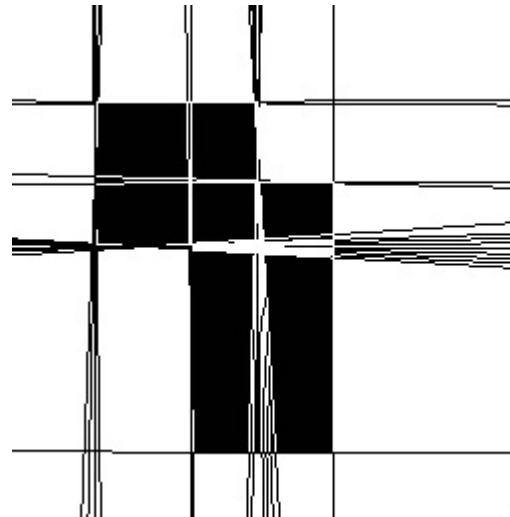
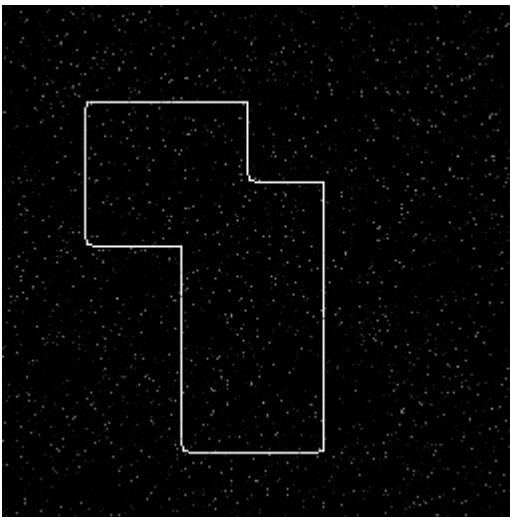


- Each  $x_i, y_i$  generates a sinusoid in  $\rho, \theta$  space

# A $\rho, \theta$ Example

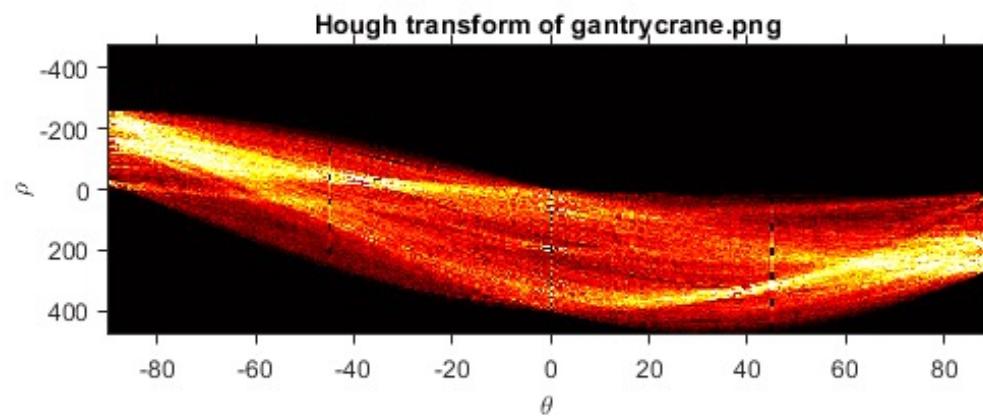


## A $\rho, \theta$ Example



*A key question is what happens to peaks in parameter space when noise or occlusion arise?*

# The Hough Transform in Matlab

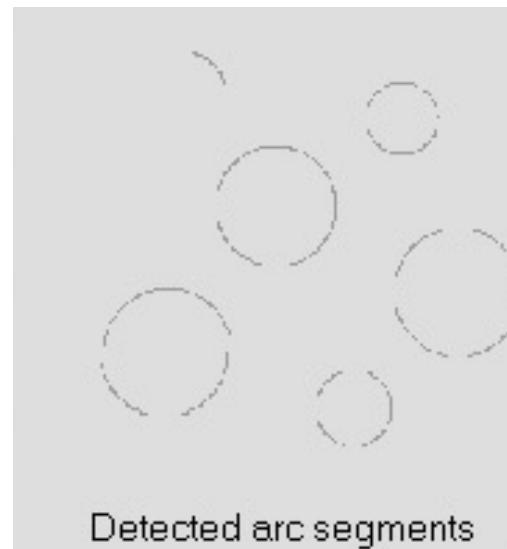
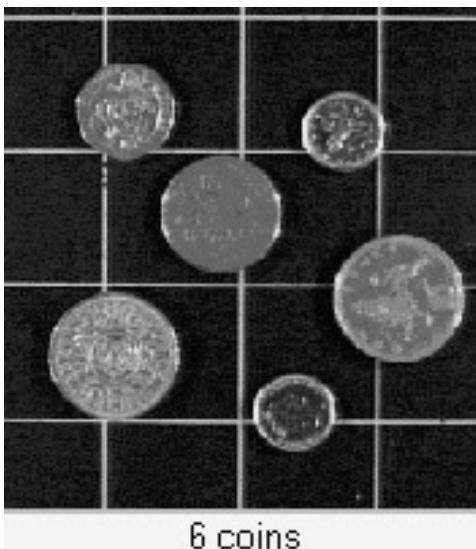


Like many implementations, Matlab's Hough allows you to specify which part of the parameter space you're interested in....

<https://uk.mathworks.com/help/images/ref/hough.html>

# Circles

- Parametric equation is  $r^2 = (x-a)^2 + (y-b)^2$ 
  - - full problem → cones in  $(a, b, r)$  space
  - - known  $r \rightarrow$  circles in  $(a,b)$  space



What about ellipses?

The Hough can in principle be applied to any parameterizable curve, though its not always practical